

# Lecture 3

## Indexing Structures for Files



# Lecture Outline

- Types of Single-level Ordered Indexes
  - Primary Indexes
  - Clustering Indexes
  - Secondary Indexes
- Multilevel Indexes
- Dynamic Multilevel Indexes Using B Trees and B+ Trees.

# Indexes as Access Paths

- Indexes are **auxiliary files** that make it more efficient to search for a record in the data file.
- The index is usually specified on **one field** of the file (although it could be specified on several fields)
- One form of an index is a file of entries  
*<field value, pointer to record/ block>*  
which is **ordered** by field value
- A **binary search** on the index yields a pointer to the file record

# Indexes as Access Paths

- The index file usually occupies considerably **less disk blocks** than the data file because its entries are much smaller.
- Indexes can be characterized as *dense* or *sparse*
  - A **dense index** has an index entry for every search key value (and hence every record) in the data file.
  - A **sparse (or non-dense) index** has index entries for only some of the search values.

# Types of Single-Level Indexes

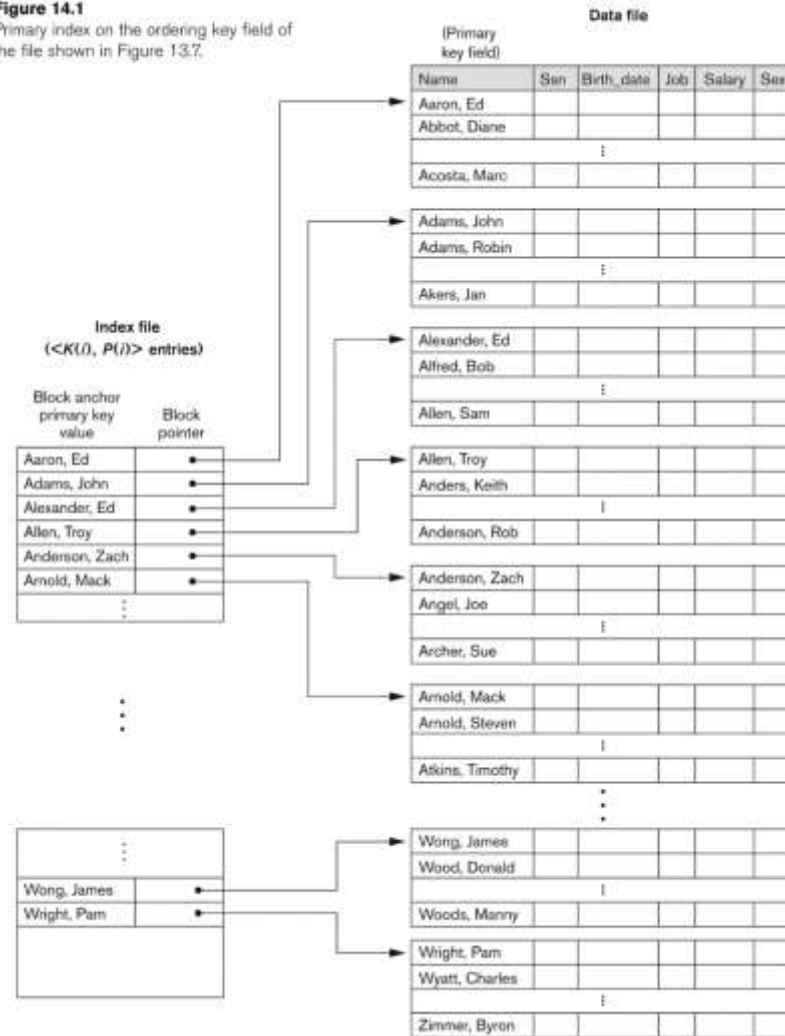
- Primary Index

- Defined on an **ordered** data file.
- The data file is ordered on a ***key field***.
- Includes one index entry for each block in the data file; the index entry has the key field value for the *first record* in the block, which is called the *block anchor*.
- A similar scheme can use the *last record* in a block.
- A primary index is a non-dense (sparse) index, since it includes an entry for each disk block of the data file and the keys of its anchor record rather than for every search value.

# Primary index on the ordering key field

**Figure 14.1**

Primary index on the ordering key field of the file shown in Figure 13.7.



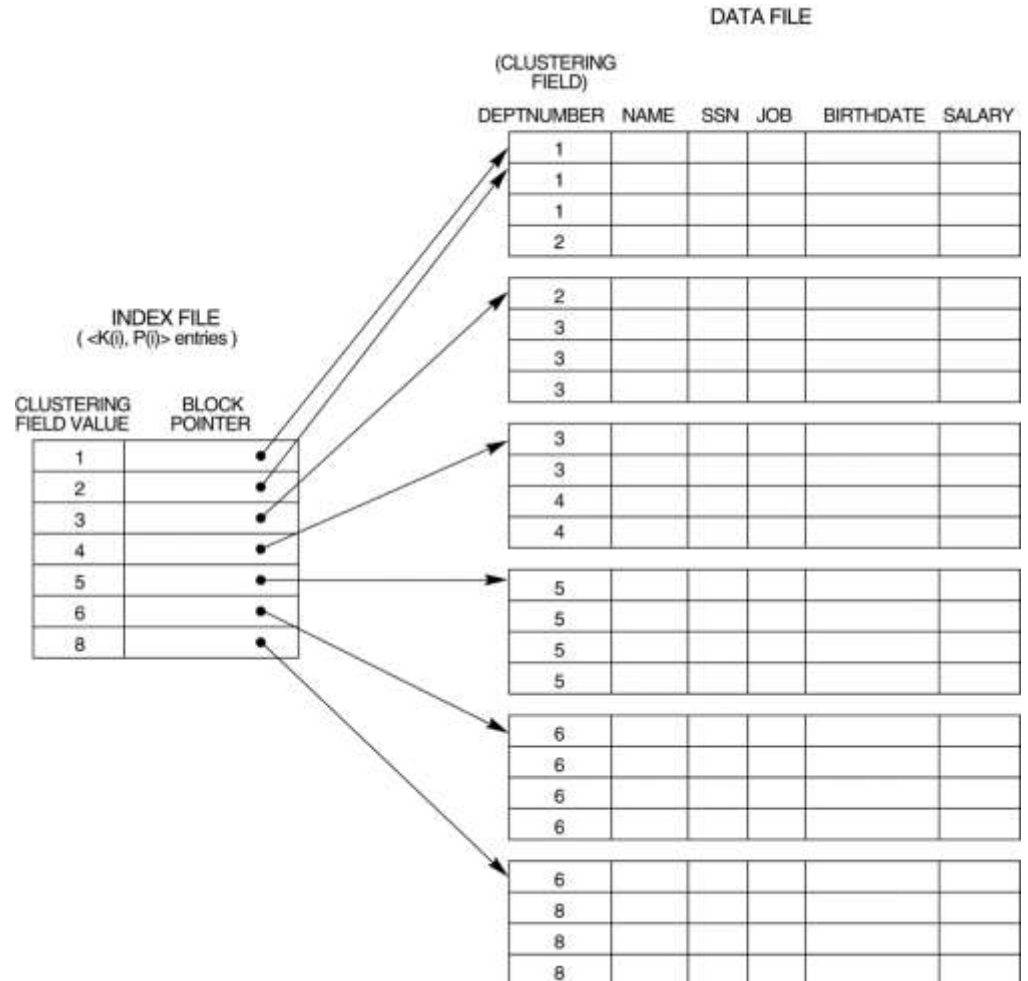
# Types of Single-Level Indexes

- Clustering Index

- Defined on an **ordered** data file
- The data file is ordered on a ***non-key field*** unlike primary index, which requires that the ordering field of the data file have a distinct value for each record.
- Includes one index entry for each distinct value of the field; the index entry points to the first data block that contains records with that field value.
- It is another example of non-dense index where Insertion and Deletion is relatively straightforward with a clustering index.

# A Clustering Index Example

A clustering index on the DEPTNUMBER (ordering non-key field of an EMPLOYEE file).

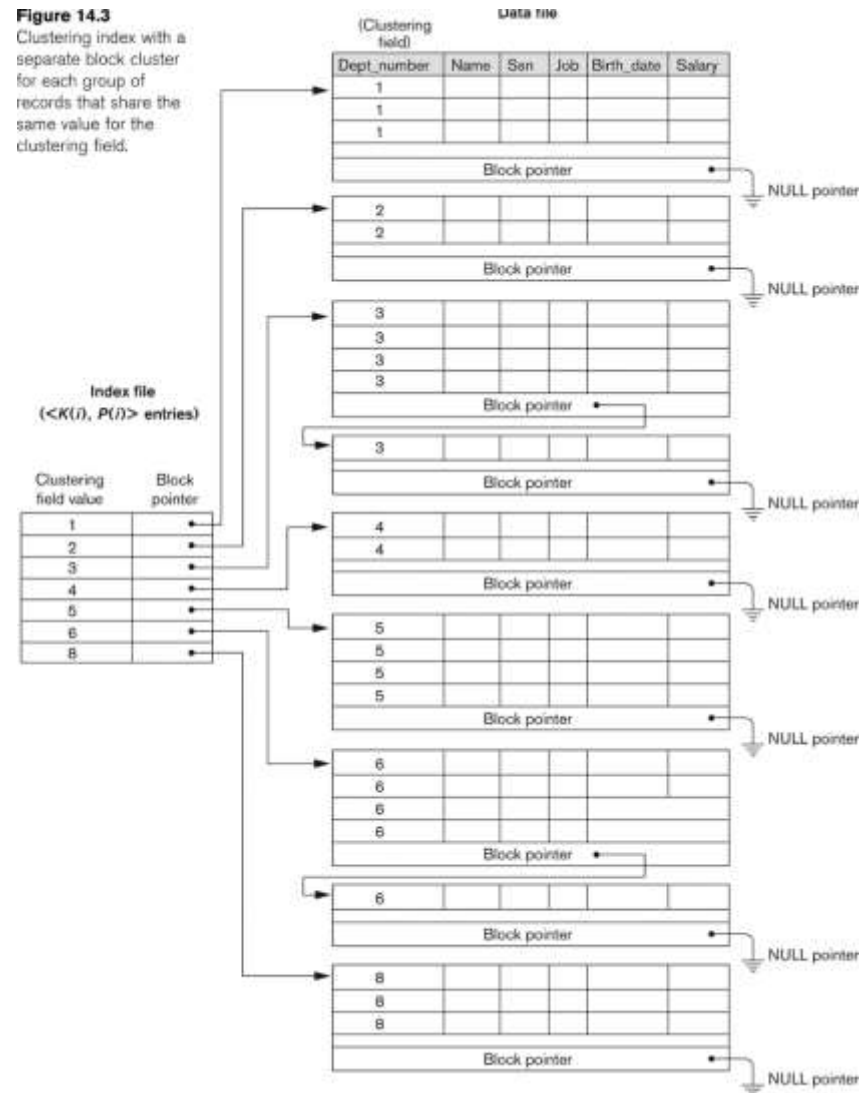




# Another Clustering Index Example

**Figure 14.3**

Clustering index with a separate block cluster for each group of records that share the same value for the clustering field.



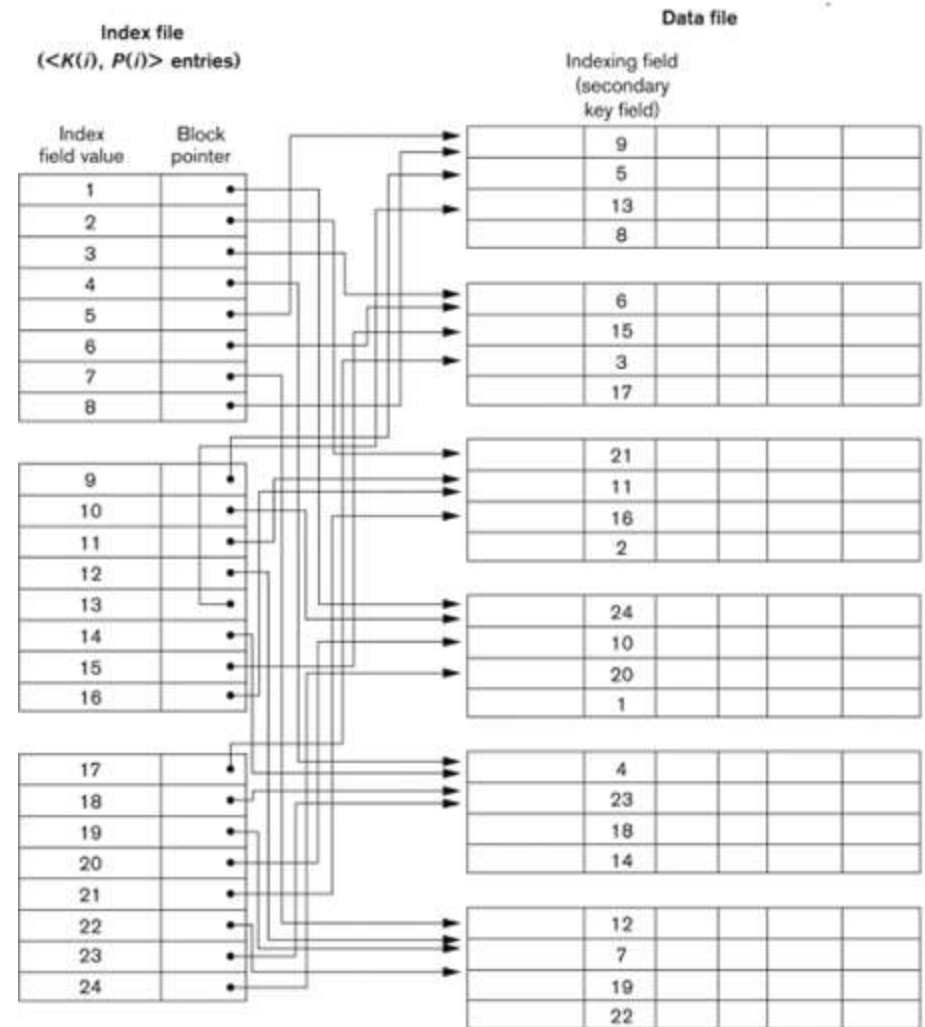
# Types of Single-Level Indexes

- Secondary Index

- A secondary index provides a **secondary means** of accessing a file for which some **primary access already exists**.
- The secondary index may be on a field which is a candidate key and has a unique value in every record, or a non-key with duplicate values.
- The index is an ordered file with two fields.
  - The first field is of the same data type as some **non-ordering field** of the data file that is an indexing field.
  - The second field is either a **block pointer or a record pointer**.
  - There can be many secondary indexes (and hence, indexing fields) for the same file.
- Includes one entry *for each record* in the data file; hence, it is a *dense index*

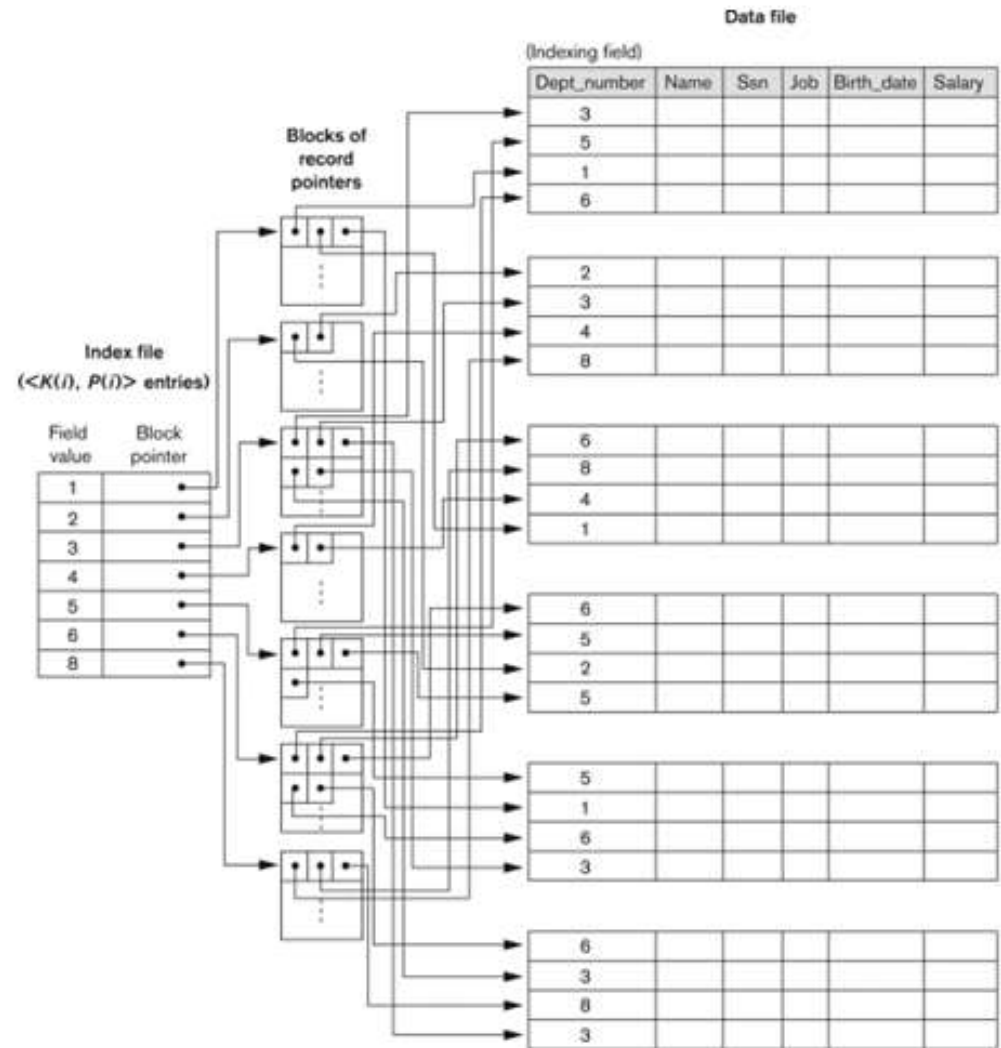
# Example of a Dense Secondary Index

A dense secondary index  
with block pointers on a  
*unordered key field*



# An Example of a Secondary Index

A dens secondary index  
with record pointers on a  
*unordered non-key field*



# Properties of Index Types

**TABLE 14.2 PROPERTIES OF INDEX TYPES**

TYPE OF INDEX	NUMBER OF (FIRST-LEVEL) INDEX ENTRIES	DENSE OR NONDENSE	BLOCK ANCHORING ON THE DATA FILE
Primary	Number of blocks in data file	Nondense	Yes
Clustering	Number of distinct index field values	Nondense	Yes/no <sup>a</sup>
Secondary (key)	Number of records in data file	Dense	No
Secondary (nonkey)	Number of records <sup>b</sup> or Number of distinct index field values <sup>c</sup>	Dense or Nondense	No

<sup>a</sup>Yes if every distinct value of the ordering field starts a new block; no otherwise.

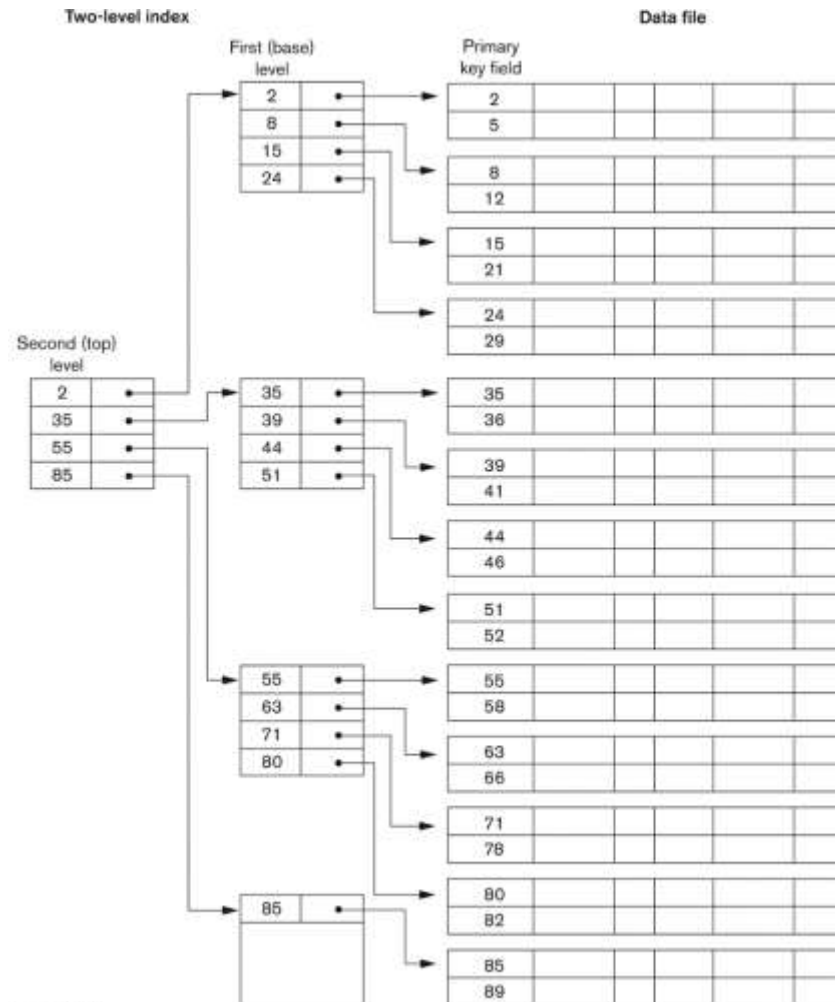
<sup>b</sup>For option 1.

<sup>c</sup>For options 2 and 3.

# Multi-Level Indexes

- Because a single-level index is an ordered file, we can create a *primary index to the index itself*;
  - In this case, the original index file is called the *first-level index* and the index to the index is called the *second-level index*.
- We can repeat the process, creating a third, fourth, ..., top level until all entries of the *top-level* fit in one disk block
- A multi-level index can be created for any type of first-level index (primary, secondary, clustering) as long as the first-level index consists of *more than one* disk block

# A Two-level Primary Index



**Figure 14.6**

A two-level primary index resembling ISAM (Index Sequential Access Method) organization.

# Multi-Level Indexes

- Such a multi-level index is a form of *search tree*
  - However, insertion and deletion of new index entries is a severe problem because every level of the index is an *ordered file*.



# A Node in a Search Tree with Pointers to Subtrees below It

- FIGURE 14.8


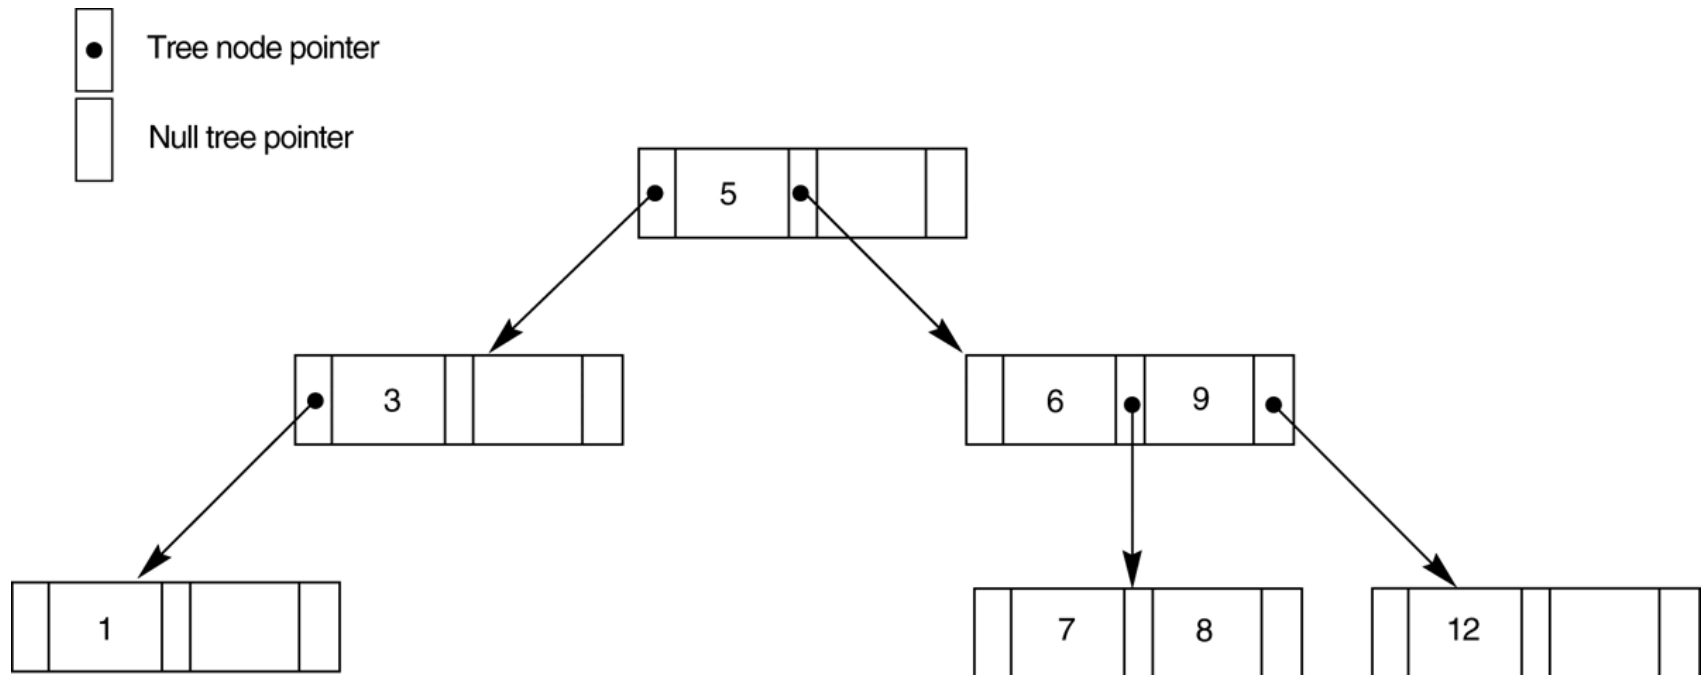


FIGURE 14.9  
A search tree of order  $p = 3$ .



# Dynamic Multilevel Indexes Using B Trees and B+ Trees

- An insertion into a node that is not full is quite efficient
  - If a node is full the insertion causes a split into two nodes
- Splitting may propagate to other tree levels
- A deletion is quite efficient if a node does not become less than half full
- If a deletion causes a node to become less than half full, it must be merged with neighboring nodes

# Difference between B-tree and B+-tree

- In a B Tree, pointers to data records exist at all levels of the tree
- In a B+ Tree, all pointers to data records exists at the leaf-level nodes
- A B+ Tree can have less levels (or higher capacity of search values) than the corresponding B Tree

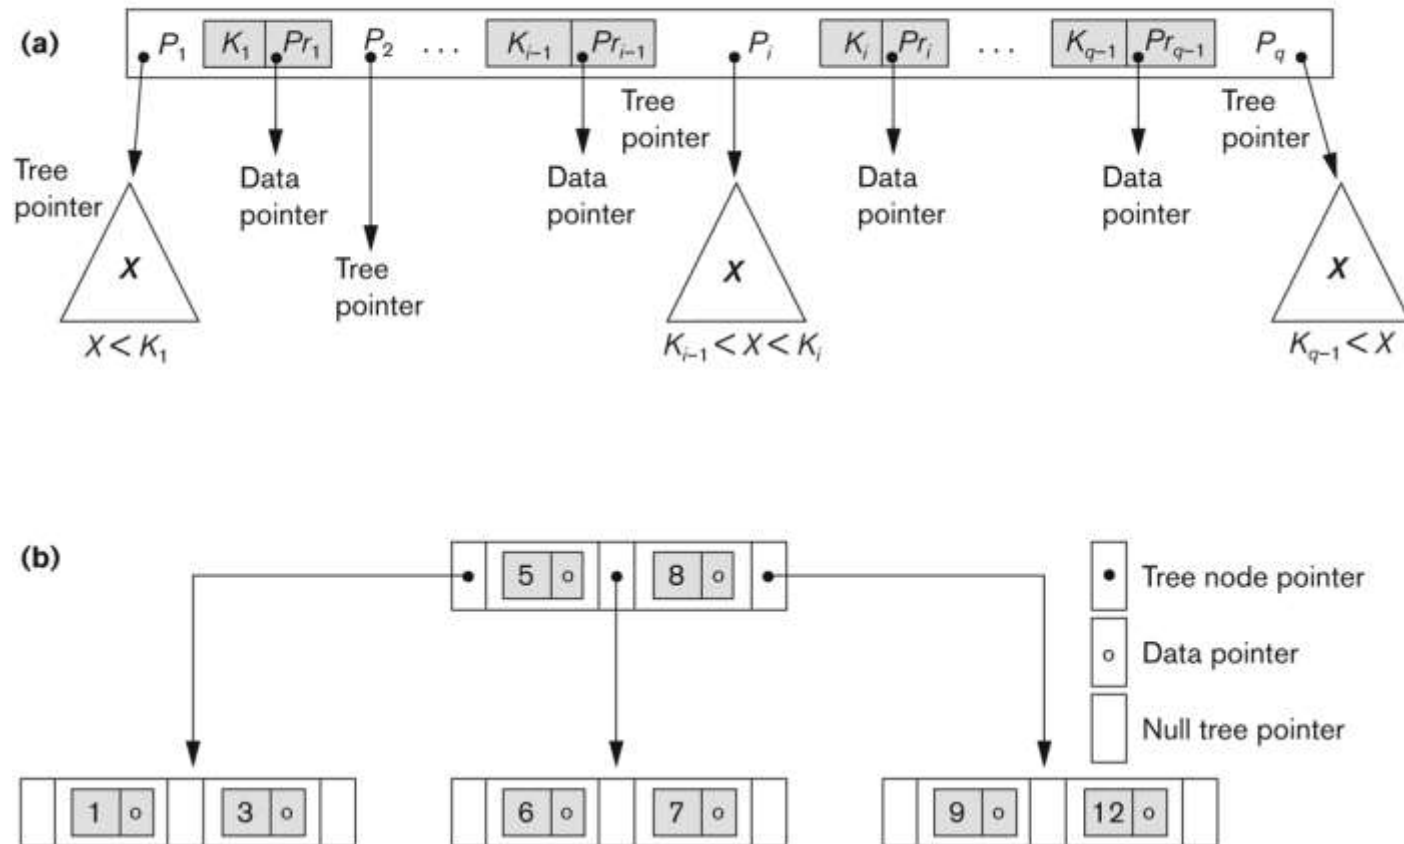
# Dynamic Multilevel Indexes Using B Trees and B+ Trees

- Most multi-level indexes use B Tree or B+ Tree data structures because of the insertion and deletion problem
  - This leaves space in each tree node (disk block) to allow for new index entries
- These data structures are variations of search trees that allow efficient insertion and deletion of new search values.
- In B Tree and B+ Tree data structures, each node corresponds to a disk block
- Each node is kept between half-full and completely full

# B Tree Structures

**Figure 14.10**

B-Tree structures. (a) A node in a B-tree with  $q - 1$  search values. (b) A B-tree of order  $p = 3$ . The values were inserted in the order 8, 5, 1, 7, 3, 12, 9, 6.



# B+ Tree Structures

- FIGURE 14.11 The nodes of a B+-tree
  - (a) Internal node of a B+-tree with  $q - 1$  search values.
  - (b) Leaf node of a B+-tree with  $q - 1$  search values and  $q - 1$  data pointers.

